

```
1 # load required packages
2 import numpy as np
3 import pandas as pd
4 pd.set_option('display.max_columns',100)
5 import matplotlib.pyplot as plt
6 %matplotlib inline
7 import seaborn as sns
8 sns.set_style('darkgrid')
9 import tensorflow as tf
10 from sklearn.metrics import confusion_matrix
11 from sklearn.preprocessing import LabelBinarizer
12 from sklearn.utils import class_weight

1 # load data
2 train_df = pd.read_csv('/content/drive/MyDrive/drugsComTrain_raw.tsv', sep = '\t')
3 test_df = pd.read_csv('/content/drive/MyDrive/drugsComTest_raw.tsv', sep = '\t')

1 # remove na values from DataFrame
2 train_df = train_df.dropna()
3 test_df = test_df.dropna()

1 # exploration
2 def explore(df):
3     print("Shape: ", df.shape, "\n")
4     print(df.dtypes, "\n")
5     print(df.head(), "\n")
6     # numeric data statistics
7     print(df.describe())
8     df.hist(figsize=(14,14), xrot=45)
9     plt.show()
10    # categorical data statistics
11    print(df.describe(include = 'object'))
12    for column in df.select_dtypes(include = 'object'):
13        if df[column].nunique() < 10:
14            sns.countplot(y = column, data = df)
15            plt.show()
16            plt.savefig('{}_dist.png'.format(column))
17
18
19 explore(train_df)
```

```
Shape: (110121, 7) \n
```

```

----- \ -----, / -----
Unnamed: 0      int64
drugName        object
condition        object
review          object
rating          float64
date            object
usefulCount     int64
dtype: object /n
      Unnamed: 0      drugName      condition \
1      95260      Guanfacine      ADHD
2      92703      Lybrel      Birth Control
3      138000      Ortho Evra      Birth Control
4      35696      Buprenorphine / naloxone      Opiate Dependence
6      165907      Levonorgestrel      Emergency Contraception



      review      rating \
1  "My son is halfway through his fourth week of ...      8.0
2  "I used to take another oral contraceptive, wh...      5.0
3  "This is my first time using any form of birth...      8.0
4  "Suboxone has completely turned my life around...      9.0
6  "He pulled out, but he cummed a bit in me. I t...      1.0

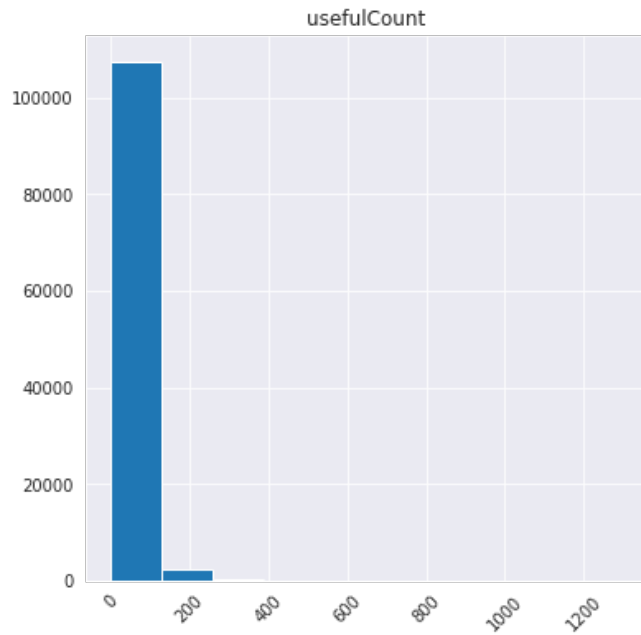


```

```

      date      usefulCount
1  April 27, 2010      192
2  December 14, 2009      17
3  November 3, 2015      10
4  November 27, 2016      37
6  March 7, 2017      5 /n
      Unnamed: 0      rating      usefulCount
count  110121.000000  110121.000000  110121.000000
mean    116603.495346    6.919924    28.008899
std      66249.766260    3.270902    38.251211
min         3.000000    1.000000    0.000000
25%      60483.000000    4.000000    6.000000
50%      117681.000000    8.000000    15.000000
75%      172110.000000   10.000000    36.000000
max     232289.000000   10.000000   1291.000000

```





	drugName	condition	review	date
count	110121	110121	110121	110121
unique	1324	32	75730	3576
top	Levonorgestrel	Birth Control	"Good"	March 1, 2016
freq	3626	28788	22	106

```
1 # obtain counts for conditions
2 condition_counts = train_df['condition'].value_counts()
3 print(condition_counts[condition_counts > 1000])
```

Birth Control	28788
Depression	9069
Pain	6145
Anxiety	5904
Acne	5588
Bipolar Disorde	4224
Insomnia	3673
Weight Loss	3609
Obesity	3568
ADHD	3383
Diabetes, Type 2	2554
Emergency Contraception	2463
High Blood Pressure	2321
Vaginal Yeast Infection	2274
Abnormal Uterine Bleeding	2096
Bowel Preparation	1859
ibromyalgia	1791
Smoking Cessation	1780
Migraine	1694
Anxiety and Stress	1663
Major Depressive Disorde	1607
Constipation	1595
Panic Disorde	1463
Chronic Pain	1455
Migraine Prevention	1413
Urinary Tract Infection	1316
Muscle Spasm	1244
Osteoarthritis	1239
Generalized Anxiety Disorde	1164
Erectile Dysfunction	1086
Opiate Dependence	1079
Irritable Bowel Syndrome	1014
Name: condition, dtype: int64	

```
1 # reduce the number of classes to anything with more than 1000 reviews
2 condition_counts_1000 = condition_counts[condition_counts > 1000]
3
4 print(len(condition_counts_1000))
5
6 counts_1000_list = list()
7 for idx, name in enumerate(condition_counts_1000.index.tolist()):
8     counts_1000_list.append(name)
9
10 train_df = train_df[train_df['condition'].isin(counts_1000_list)]
11 test_df = test_df[test_df['condition'].isin(counts_1000_list)]
12
13 print(train_df['review'].shape)

32
(110121,)


1 print(counts_1000_list)

['Birth Control', 'Depression', 'Pain', 'Anxiety', 'Acne', 'Bipolar Disorde',
```

```

1 # visualizations
2 fig, ax = plt.subplots(1, 1)
3 ax = sns.countplot(x = 'condition', data = train_df)
4 ax.set_title('Distribution of Conditions')
5 ax.set_xticklabels(ax.get_xticks(), rotation = 45)
6 fig.show()
7 fig.savefig('condition_dist.png')

```



```

1 # set data and labels
2 x_train = train_df['review']
3 x_test = test_df['review']
4 y_train = train_df['condition']
5 y_test = test_df['condition']

1 # labels to one hot encoded
2 def prepare_targets(y_train, y_test):
3     one_hot = LabelBinarizer()
4     one_hot.fit(y_train)
5     y_train = np.argmax(one_hot.transform(y_train), axis = 1)
6     y_test_one_hot = one_hot.transform(y_test)
7     y_test = np.argmax(one_hot.transform(y_test), axis = 1)
8     y_test_rev = one_hot.inverse_transform(y_test_one_hot)
9     return y_train, y_test, y_test_one_hot, y_test_rev
10
11 y_train, y_test, y_test_one_hot, y_test_rev = prepare_targets(y_train, y_test)

```

```
1 # create weights for classes
2 class_weights = class_weight.compute_class_weight('balanced', np.unique(y_train
3
4 weights = dict(enumerate(class_weights))
5 print(weights)
```

```
72276825302986, 1: 1.6418326574427482, 2: 0.6158341535433071, 3: 0.58287284044
```

```
1 # baseline model; all predictions birth control
2 count_bc = condition_counts['Birth Control']
3 base_acc = count_bc/len(y_train)
4 print(base_acc)
```

```
0.26142152722913886
```

```
1 # vectorize text
2 vocab_size = 1500
3 review_len_max = 200
4 encoder = tf.keras.layers.experimental.preprocessing.TextVectorization(
5     max_tokens = vocab_size,
6     output_sequence_length = review_len_max)
7
8 # develop vocabulary
9 encoder.adapt(x_train.values)
10
11 # vectorize text
12 x_train = encoder(x_train)
13 x_test = encoder(x_test)
```

```
1 print('Training input shape: ', x_train.shape)
2 print(len(x_train))
3 print(len(x_train[0]))
4 print(x_train[0].shape)
5 print('Test input shape: ', x_test.shape)
6 print(len(x_test))
7 print(len(x_test[0]))
8 print(x_test[0].shape)
```

```
Training input shape: (110121, 200)
110121
200
(200,)
Test input shape: (36827, 200)
36827
200
(200,)
```

```
1 print(y_train.shape)

(110121,)
```



```

1 # classifiying with tf.keras RNN
2 model = tf.keras.models.Sequential([
3     tf.keras.layers.Embedding(input_dim = vocab_size + 1, output_dim = 16),
4     tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(units = 16, dropout = 0.5)),
5     tf.keras.layers.Dense(32, activation = 'softmax')
6 ])
7 )
8
9 # compile model
10 model.compile(loss = tf.keras.losses.SparseCategoricalCrossentropy(),
11               optimizer = 'Adam',
12               metrics = ['acc'])
13
14 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 16)	24016
bidirectional (Bidirectional)	(None, 32)	4224
dense (Dense)	(None, 32)	1056
Total params: 29,296		
Trainable params: 29,296		
Non-trainable params: 0		

```

1 # test model
2 history = model.fit(x = x_train, y = y_train,
3                     epochs = 50,
4                     batch_size = 50,
5                     validation_split = 0.2,
6                     class_weight = weights,
7 )

```

```

Epoch 1/50
1762/1762 [=====] - 68s 20ms/step - loss: 3.2001 - a
Epoch 2/50
1762/1762 [=====] - 34s 19ms/step - loss: 2.6470 - a
Epoch 3/50
1762/1762 [=====] - 34s 19ms/step - loss: 2.3106 - a
Epoch 4/50
1762/1762 [=====] - 34s 20ms/step - loss: 2.0920 - a
Epoch 5/50

```

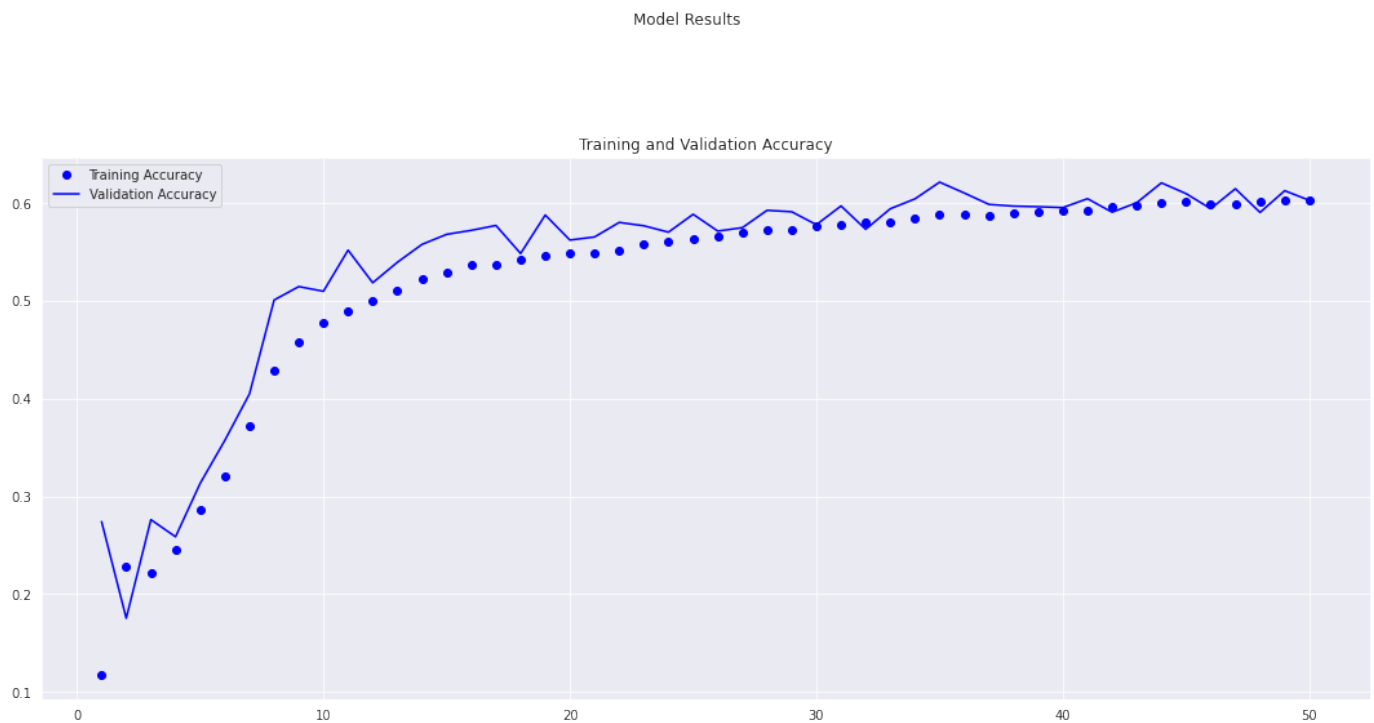
```
Epoch 5/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.9465 - a
Epoch 6/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.8344 - a
Epoch 7/50
1762/1762 [=====] - 33s 19ms/step - loss: 1.7504 - a
Epoch 8/50
1762/1762 [=====] - 30s 17ms/step - loss: 1.6362 - a
Epoch 9/50
1762/1762 [=====] - 33s 19ms/step - loss: 1.5683 - a
Epoch 10/50
1762/1762 [=====] - 35s 20ms/step - loss: 1.4990 - a
Epoch 11/50
1762/1762 [=====] - 33s 19ms/step - loss: 1.4578 - a
Epoch 12/50
1762/1762 [=====] - 31s 17ms/step - loss: 1.4281 - a
Epoch 13/50
1762/1762 [=====] - 31s 18ms/step - loss: 1.3844 - a
Epoch 14/50
1762/1762 [=====] - 35s 20ms/step - loss: 1.3514 - a
Epoch 15/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.3351 - a
Epoch 16/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.3112 - a
Epoch 17/50
1762/1762 [=====] - 32s 18ms/step - loss: 1.2995 - a
Epoch 18/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.2925 - a
Epoch 19/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.2697 - a
Epoch 20/50
1762/1762 [=====] - 35s 20ms/step - loss: 1.2583 - a
Epoch 21/50
1762/1762 [=====] - 34s 19ms/step - loss: 1.2514 - a
Epoch 22/50
1762/1762 [=====] - 35s 20ms/step - loss: 1.2458 - a
Epoch 23/50
1762/1762 [=====] - 35s 20ms/step - loss: 1.2107 - a
Epoch 24/50
1762/1762 [=====] - 33s 19ms/step - loss: 1.1952 - a
Epoch 25/50
1762/1762 [=====] - 31s 17ms/step - loss: 1.2085 - a
Epoch 26/50
1762/1762 [=====] - 31s 18ms/step - loss: 1.1817 - a
Epoch 27/50
1762/1762 [=====] - 32s 18ms/step - loss: 1.1740 - a
Epoch 28/50
1762/1762 [=====] - 33s 19ms/step - loss: 1.1762 - a
Epoch 29/50
1762/1762 [=====] - 32s 18ms/step - loss: 1.1684 - a
Epoch 30/50
```


```

1762/1762 [-----] 24s 10ms/step loss: 1.1671
1 # save model
2 model.save('drug_review_rnn_50.h5')

1 # model results
2 acc = history.history['acc']
3 val_acc = history.history['val_acc']
4 loss = history.history['loss']
5 val_loss = history.history['val_loss']
6 epochs = range(1, len(acc) + 1)
7
8 fig, (ax0, ax1) = plt.subplots(2, 1, figsize = (15,15))
9 fig.suptitle('Model Results')
10 ax0.plot(epochs, acc, 'bo', label = 'Training Accuracy')
11 ax0.plot(epochs, val_acc, 'b', label = 'Validation Accuracy')
12 ax0.set_title('Training and Validation Accuracy')
13 ax0.legend()
14
15 ax1.plot(epochs, loss, 'bo', label = 'Training Loss')
16 ax1.plot(epochs, val_loss, 'b', label = 'Validation Loss')
17 ax1.set_title('Training and Validation Loss')
18 ax1.legend()
19 fig.tight_layout(rect=[0, 0.03, 1, 0.9])
20
21 fig.savefig('model_results.png')

```





```
1 # model performance
2 predictions = model.predict(x_test)
```

```

1 print(predictions[0])
2 print(y_test[0])
3 print(y_test_one_hot[0])
4 print(y_test_rev[0])

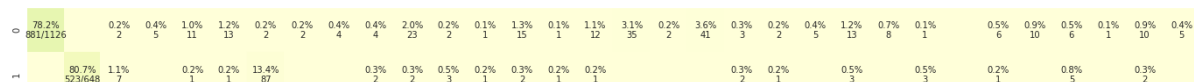
[4.1943714e-03 5.3141492e-05 1.1367062e-04 9.5583566e-02 2.2228974e-01
 1.9638129e-02 1.8428185e-03 7.2395586e-08 2.5014702e-05 5.0075510e-06
 2.4951330e-01 7.3006345e-06 3.9542447e-06 1.4613361e-06 1.5919979e-01
 1.8198247e-04 1.3924406e-03 4.0000530e-05 2.0378804e-01 6.9832919e-07
 7.5372423e-05 1.2797008e-06 2.0408104e-04 6.0204049e-05 4.9668624e-06
 5.3810249e-06 4.0609650e-02 6.3067378e-04 3.0860517e-06 4.8659804e-06
 1.6989521e-04 3.5602206e-04]
10
[0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
Depression

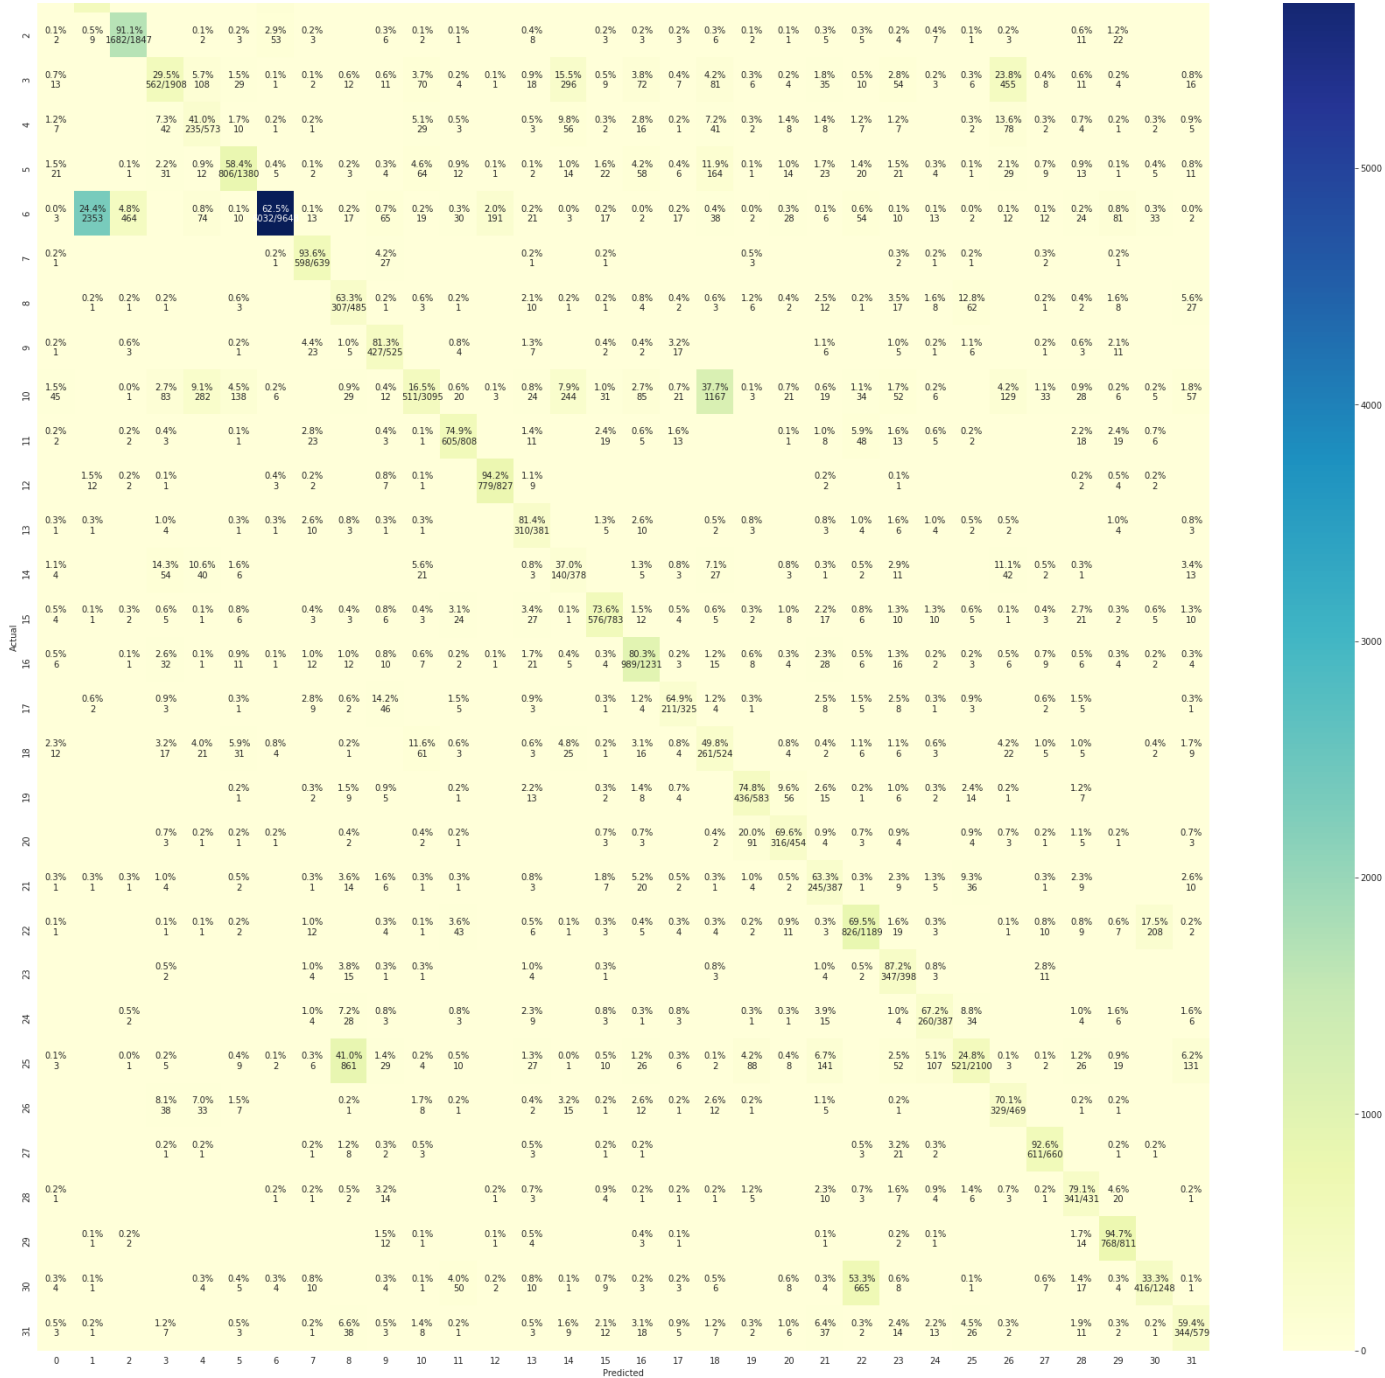
```

```

1 # build confusion matrix
2 def plot_cm(y_true, y_pred, figsize = (30, 30)):
3     cm = confusion_matrix(y_true, y_pred, labels = np.unique(y_true))
4     cm_sum = np.sum(cm, axis = 1, keepdims = True)
5     cm_perc = cm / cm_sum.astype(float) * 100
6     annot = np.empty_like(cm).astype(str)
7     nrows, ncols = cm.shape
8     for i in range(nrows):
9         for j in range(ncols):
10             c = cm[i, j]
11             p = cm_perc[i, j]
12             if i == j:
13                 s = cm_sum[i]
14                 annot[i, j] = '%.1f%%\n%d/%d' % (p, c, s)
15             elif c == 0:
16                 annot[i, j] = ''
17             else:
18                 annot[i, j] = '%.1f%%\n%d' % (p, c)
19     cm = pd.DataFrame(cm, index = np.unique(y_true), columns = np.unique(y_true))
20     cm.index.name = 'Actual'
21     cm.columns.name = 'Predicted'
22     fig, ax = plt.subplots(figsize = figsize)
23     sns.heatmap(cm, cmap= "YlGnBu", annot = annot, fmt = '', ax = ax)
24     fig.savefig('confusion_matrix.png')
25
26 plot_cm(y_test_one_hot.argmax(axis = 1), predictions.argmax(axis = 1))

```





1

 10s completed at 9:36 PM